

The 3-SAT Decision Problem

Towards a Parallel Search Implementation

Team Satisfaction

Christopher Wood, Eitan Romanoff, Ankur Bajoria

May 7, 2013

Agenda

- 1 Problem Statement
- 2 Exhaustive Search Algorithm
- 3 Advanced Techniques and Heuristics for Parallel SAT Algorithms
- 4 Parallel Program Design and Demonstration
- 5 Exhaustive Performance Metrics
- 6 Lessons Learned and Future Work
- 7 Questions

3-SAT Problem

Boolean satisfiability is an *NP*-complete decision problem defined as:

$$SAT : \phi \rightarrow \{YES, NO\}$$

Input: 3-CNF Boolean formula ϕ_n on n variables.

$$\phi_3 = (x_1 \vee x_2 \vee \neg x_3) \wedge (\neg x_1 \vee x_2 \vee x_3) \wedge (\neg x_1 \vee x_2 \vee \neg x_3)$$

Output: *YES* if there exists a truth assignment to the variables in ϕ_n such that it evaluates to true, *NO* otherwise.

$$\phi_n \text{ is satisfiable} \Leftrightarrow SAT(\phi_n) = YES$$

Exhaustive Search for 3-SAT

Input: 3-CNF formula ϕ_n on n variables, **Output:** YES or NO

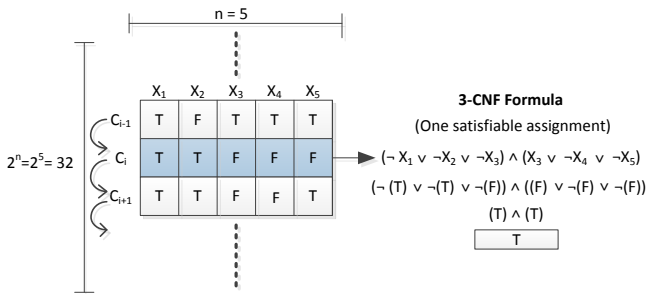
```
1:  $C \leftarrow FALSE^n$  (array of  $n$  False values, the initial configuration)
2: for  $i = 0 \rightarrow 2^n - 1$  do
3:    $SAT \leftarrow TRUE$ 
4:   for all  $clause \in \phi_n$  do
5:     if  $evaluate(clause, C) = FALSE$  then
6:        $SAT \leftarrow FALSE$ 
7:     end if
8:   end for
9:   if  $SAT = TRUE$  then return YES
10:   $C \leftarrow nextConfig(C)$ 
11: end for
12: return NO
```

Exhaustive Search for 3-SAT - An *Exhaustive* Example!

Input: $\phi_5 = (\neg X_1 \vee \neg X_2 \vee \neg X_3) \wedge (X_3 \vee \neg X_4 \vee \neg X_5)$

Output: Yes

Note: no early termination once a satisfiable solution is found



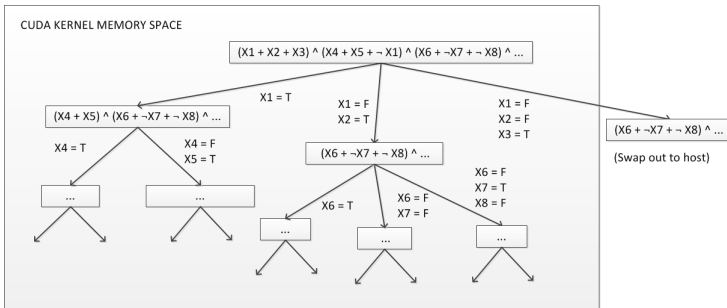
Evaluating a Clause

- Evaluating a clause depends on how ϕ_n and the variable truth assignments are stored.
 - `boolean[] variables` for variable assignments and `Literal[][3] formula` for ϕ_n .
 - A `Literal` has a variable ID and negated flag

```
for (int c = 0; c < numClauses; c++) {
    boolean clauseValue = false;
    for (int l = 0; l < 3 && clauseValue == false; l++) {
        if (formula[c][l].negated == true && !variables[formula[c][l].id])
            clauseValue = true;
        else if (formula[c][l].negated == false && variables[formula[c][l].id])
            clauseValue = true;
    }
    // Check the value of the clause now...
}
```


Recursive Divide and Conquer with CUDA

- Idea: Recursively search the variable assignment space
- Approach: A formulas is a “stack” of clauses that are reduced or popped off during assignment
- Implementation: Master-worker pattern for distributing formulas to individual CUDA kernels




Advanced Clause Learning and Sharing with ManySAT

- Idea: Exploit heuristics for particular types of formulas
- Approach: Run multiple sequential solvers in parallel and reduce the result together
- Implementation: Each solver uses different restart policies, literal selection strategies, and shared conflict-driven clause learning



Intelligent Literal Decisions and Advanced Data Structures for DPLL

- Idea: Enhance literal selection to more effectively traverse the search space
- Approach: Use trie data structures to build “guided paths” that indicate whether a branch has been visited or not
- Implementation: Distribute guided paths across different worker processes using a master-worker pattern



Parallel Program Characteristics

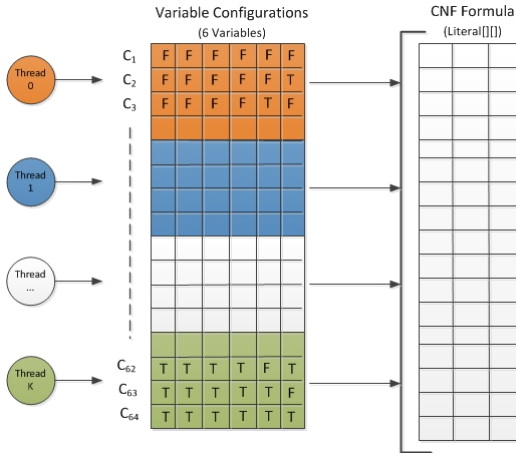
Our design goals included:

- Evenly divide the computation among different threads
- Minimize (or remove) conflicts for shared variables


Our parallel design strategy:

- *Result parallelism* for exhaustive program and *agenda parallelism* for decision program
- Split the evaluation of each variable configuration among every thread
- Reduce the final result into the main thread

Computation Partition Strategy



Thread Synchronization







Action!





Demo time!

$$\phi_5 = (\neg X_1 \vee \neg X_2 \vee \neg X_3) \wedge (X_3 \vee \neg X_4 \vee \neg X_5)$$

Speedup Metrics (Exhaustive)



Sizeup Metrics (Exhaustive)



Performance Observations

- We achieved *superlinear* speedups and sizeups when varying the number of variables
- Our parallel programs achieve better performance when the number of variables was varied:
 - The problem size $N = f(N_v, N_c) = 2^{N_v} \times N_c$

Lessons Learned

- If the problem size is a function of *multiple* variables, experiments should only change one of such variables to gather valid performance data
- Exploiting the cache and JVM can yield *extremely* good speedup and sizeup efficiencies

Future Work

- Implement more advanced heuristics for literal selection
- Strive for wider splits of the configuration search space tree among multiple processes
- Experiment with different data structures to see what's the most optimal

Questions?

Fire away!